
15-418/618 Project Report

MercuryJson: Multi-Threaded JSON Parsing with SIMD

Zecong Hu

Carnegie Mellon University
zeconghu@andrew.cmu.edu

Da Yin

Carnegie Mellon University
dayin@andrew.cmu.edu

1 Summary

We have implemented a parallel JSON parser with SIMD and multi-threading. Our implementation using 4 threads achieves on average $1.29\times$ speedup on large documents compared to a very strong baseline, and achieves comparable performance on small documents.

2 Background

2.1 JSON Format

JSON (JavaScript Object Notation) [1, 3] is a lightweight data-interchange format based on a subset of the JavaScript Programming Language, Standard ECMA-262 3rd Edition - December 1999.

The EBNF (Extended Backus-Naur form) of JSON grammar is described in figure 1. JSON is built on two structures, namely objects and arrays. Objects are collections of key-value pairs and arrays are ordered lists of values. Keys are defined as strings composed of a sequence of zero or more Unicode characters, wrapped in double quotes, using backslash escapes. Values could be strings, numbers (integer or floating-point), keywords (`true`, `false`, `null`), or objects or arrays as described above.

2.2 JSON Parsing

There are many existing JSON parsers used in various programming languages, most of which are based on finite state machines (FSM). The state-of-the-art parser so far is *simdjson* which utilizes SIMD instructions to accelerate the parsing process.

simdjson [4] divides the traditional FSM-based parsing into two stages. This two-stage method was first proposed in *Mison* [6], a fast JSON parser by Microsoft, based on structure speculation. *Mison* uses large datasets to learn to predict structural character positions which greatly eliminates branch prediction. While they can achieve over 2 GB/s parsing speed as well, the speed across different datasets varies.

In *simdjson*, by elegantly devising vector operations and using loop-unrolling tricks to avoid frequent branch predictions, it can reach over 2 GB/s by exploiting parallelism in first stage which is independent of datasets. The first stage includes finding structural characters such as brackets and commas after picking out escape backslash characters and identifying literal values. Then, the parser will calculate the offsets between structural characters for use in second stage. Although the second stage for both *Mison* and *simdjson* are sequential, the usage of offset between structural characters greatly benefits parsing speed as it allows skipping over most content characters. *simdjson* also utilizes a tape-based storage for fast creation of JSON structures, at the cost of slow random access.

```

json           = value

value          = string | number | object | array | "true" | "false" | "null"

object         = "{" , object-elements , "}"
object-elements = string , ":" , value , [ "," , object-elements ]

array         = "[" , array-elements , "]"
array-elements = value , [ "," , array-elements ]

string         = '"' , { character } , '"'
character      = "\" , ( '"' | "\" | "/" | "b" | "f" | "n" | "r" | "t" | "u" , digit ,
                digit , digit ) | unicode

digit          = "0" | digit-1-9
digit-1-9      = "1" | "2" | "3" | "4" | "5" | "6" | "7" | "8" | "9"
number         = [ "-" ] , ( "0" | digit-1-9 , { digit } ) , [ "." , { digit } ] , [ ( "e" |
                "E" ) , ( "+" | "-" ) , { digit } ]

```

Figure 1: Extended Backus-Naur form of JSON grammar. The term “unicode” is defined as any Unicode character except those that must be escaped. Please refer to the JSON standards for details.

Our work is mainly based on *simdjson*. We also base our parser on the two-stage parsing procedure, and introduce multi-threading to further accelerate the second stage.

2.3 General Parsing Algorithms

Our work is a small piece in the broad field of parsing. Although very different in implementation, our methods are in essence extensions built upon the fundamental parsing algorithms, namely **recursive descent** and **shift-reduce**.

Recursive descent parsers are top-down parsers built from a set of mutually recursive procedures (or a non-recursive equivalent) where each such procedure implements one of the nonterminals of the grammar [2]. Recursive descent starts with the start symbol, and recursively expands each nonterminal symbol using production rules. Non-backtracking recursive descent parsers support $LL(k)$ grammars, which is a subset of unambiguous context-free grammars.

Shift-reduce parsers, on the other hand, are bottom-up parsers that incrementally builds the parse tree from bottom up. Shift-reduce parsers scan text from left to right, applying shift steps and reduces steps alternatively. A shift step consumes an input token, creates a single-node parse tree, and pushes it onto the parse stack. A reduce step applies a production rule to several parse trees off the stack top, replacing them with a joined tree with a new root symbol. Shift-reduce parsers support $LR(k)$ grammars, also a subset of unambiguous context-free grammars.

The JSON grammar is classified as both $LL(1)$ and $LR(1)$, thus both parsing algorithms are applicable to parsing JSON.

3 Approach

As mentioned, our approach uses a two-stage framework similar to that of *simdjson*. Before introducing these two stages, we first define **structural characters** as the set of all square and curly brackets, colons, commas, and the first character of literal values (always quotes (")), numerical values (minus sign (-) or digits 0 through 9), and keyword values (t for true, f for false, n for null). Among these characters, brackets, colons, and commas describe the JSON structure, while others indicate the beginning of a JSON value. If the document is valid, then structural characters alone are sufficient to define the JSON structure, allowing us to skip over most characters in the second stage.

The first stage of parsing is to find the positions of structural characters. We reimplemented SIMD parsing algorithm in *simdjson*. The original input is divided into groups size of V , where V is the maximum data width that the processor supports divided by 8 (each character takes 8 bits). This could be 128, 256, or 512 bits depending on processor architecture. In our experiments, we use 256-bit AVX2 as default and 128-bit as fallback.

1. Compute the escape mask. Backslash characters (\) escape their following characters, and the aim of this step is to find out which characters need to be escaped. Continuous backslashes makes this computation a bit complicated. The process is shown as below. Note that our implementation is a little different from *simdjson* — while *simdjson* only extracts ending escape positions, we extract all positions for later use.

2. Compute the literal mask which indicates which positions are inside a literal (string) value. By finding unescaped quotes and using carry-less multiplication `clmul`, we can compute the literal mask in only several cycles.

3. Compute white space and structural masks. Here we use the `vpshufb` instruction to do mapping of white space characters and ordinary structural characters (brackets, colons, and commas). This trick is the same as in `simdjson`. `vpshufb` can only handle 16 bits mapping so we use a two-way mapping (one for higher 8 bits and one for lower 8 bits), and then bitwise-AND the mapped results.

```
{ "Nam[{" : [ 116,"\\\\" , 234,"true", false ], "t":"\\" }]: input
_1_____1_____1_____1_____1_1_1_____: Q (Quotes)
_1111111111_1111_1111_1111_____11_1111____: R (Literal mask)
1_____1_1_____1_____1_____1_____1_1_____1: S (Structural)
_1_____1_1_____1_1_____1_____1_1_____1: W (Whitespace)
// structural and whitespace mask
111_____1111_11_1111___11_____111_111_111___111: P = S | W
// add quotes to the structural mask
1_1_____11_1_____11_____1_1_____1_1_____11_1_111___1_1: S = S | Q
// obtain candidates: structural characters must be preceded by whitespace
// or another structural character
1111_11111_11_1111_111_111_111_11111_111_11: P = P << 1
```

```
// eliminate whitespace and quoted characters
1_1_1_1_1_1_1_1_1_1_1_1_1_1_1_1_1_1_1_1_1_1:P=&~W&~R
// merge structural characters
1_1_1_1_1_1_1_1_1_1_1_1_1_1_1_1_1_1_1_1_1_1:S=S|P
// eliminate ending quotes
1_1_1_1_1_1_1_1_1_1_1_1_1_1_1_1_1_1_1_1_1_1:S&~(Q&~R)
```

5. Extract positions of structural characters from the mask into a continuous indices array. This can be implemented using `tzcnt` and `blsr` instructions.

Note that computation for each group depends on escape mask and literal mask results from the previous group.

3.2 Representation of the JSON Structure

We experiment with different representations of the JSON document structure. Here we summarize each representation:

- **JSON Value Nodes** (§3.2.1): The JSON document is stored as a collection of nodes, where each node encapsulates a JSON value. Key–value pairs in objects, and values in arrays are stored as a linked list.
- **Tape-based Storage** (§3.2.2): The JSON document is stored as a *tape*, i.e. a “flattened” representation of the JSON document, obtained by traversing the JSON structure in preorder. This is the representation used in *simdjson*.

We must emphasize that these representations are not optimized for random access, but for creation and traversal. For instance, one common use case for JSON documents might be to access a specific value nested under specific keys. For either representation, performance could degenerate to $O(N)$ for a document of length N . Common solutions to these use cases include using map data structures (e.g. `std::map`) for objects and variable-size lists (e.g. `std::vector`) for arrays, but either incurs unbearable overhead. We argue that while these may be common use cases, they are not typical of users seeking to parse JSON documents at gigabytes per second.

3.2.1 JSON Value Nodes

In this representation, each JSON value is stored as an instance of `JSONValue` type. The instance contains two fields: an integer field describing the value type, and a union field storing the actual value. For keyword types, the union field is empty; for number types, the union field stores a 64-bit integer or floating-point value; for string values, the union field stores a pointer to the first character of a null-terminated string; for array and object values, the union field stores a pointer to a `JSONArray` or `JSONObject` instance. `JSONArray` and `JSONObject` are linked lists, the former containing pointer to a `JSONValue` instance, the latter containing pairs of pointers to a string and a `JSONValue` instance.

Since this representation requires dynamic allocation of many small objects, to prevent memory fragmentation and speedup allocation, we used a blocked memory allocator that first allocates a fixed-sized, aligned block of memory, and constructs JSON values on this block of memory. When the block has exhausted, a new block is allocated. With properly chosen block sizes, the allocator requires on average less than two memory allocations, and results in very few branch misses.

The advantage of this representation is its flexibility, since each node on its own represents a parse tree, which could potentially be partially parsed. This makes it parser-agnostic — applicable to both top-down and bottom-up parsers. The disadvantage is excessive use of space and memory access overhead due to storing pointers.

3.2.2 Tape-based Storage

The tape-based storage is basically a condensed representation of the input string. JSON values are stored on the tape in the same order as they are parsed. Each element of the tape is a 64-bit unsigned integer, representing either a JSON value or a structural character. The higher 4 bits contain the value type, and the remaining 60 bits store the payload. Values of each type are represented as follows:

Document	Tape		Numerals	Literals
{	0: {	11	...	0- 7: ..name..
"n ame ": [1: string	2	4: 116	8-15:
116,	2: [8	...	16-23: \\.....
"\\",	3: number	4	8: 234	24-31:tru
234,	4: string	16	...	32-39: e.....
"true",	5: number	8		40-47:t...
false	6: string	29		48-54: .\"....
],	7: false			
"t": "\\\""	8:]	2		
}	9: string	44		
	10: string	49		
	11: }	0		

Figure 2: Example of tape-based storage. Columns from left to right contain: the original JSON document, the parsed tape storage, the numerals array, the string literals array (‘.’ represents null character ‘\0’). Each line in the tape column contains: tape offset, value type, payload.

- **Keywords** (true, false, null): 1 element, payload is empty.
- **Number**: 1 element, payload stores the offset into a `numerals` array.
- **String**: 1 element, payload stores the offset into a `literals` array.
- **Array**: 2 elements, corresponding to opening and closing square brackets, each storing in payload the tape offset of the other element. All values (including objects and arrays) in the array are written in between the two elements.
- **Object**: 2 elements, corresponding to opening and closing curly brackets, each storing in payload the tape offset of the other element. All key-value pairs (including object and array values) in the object are written alternately in between the two elements.

Our tape design is similar to the one used in *simdjson*, with the only difference being that we store numbers off-tape. This gives us the guarantee that the tape length never exceeds the number of structural characters. Figure 2 shows an example of tape-based storage.

3.3 Parsing the JSON Structure

We experimented with several different algorithms for structured parsing, each having different capabilities. Here we summarize each algorithm:

- **Recursive Descent** (§3.3.1): Our baseline approach.
- **Finite State Machine** (§3.3.2): The approach used in *simdjson*. By manually maintaining the program stack, recursive descent can be simulated as a finite state machine (FSM) with lower overhead.
- **Shift-Reduce** (§3.3.3): A bottom-up parsing algorithm that can be parallelized using multi-threading. Structural characters are shifted onto a stack, and elements from stack top are reduced according to production rules. Nodes representation must be used.
- **Multi-Thread FSM** (§3.3.4): Our final and fastest approach. This algorithm supports parsing of incomplete structure by including “unknown states”, where the parser can be in several possible states. Its difference with single-threaded state machine is similar to that between an NFA and a DFA. Although nodes representation is also applicable, we opt for tape-based storage for its efficiency.

3.3.1 Recursive Descent

For this method, we used the most straightforward implementation: a set of mutually recursive procedures for each nonterminal, i.e. JSON values, arrays, and objects. Implementation is based on a simplified version of the JSON EBNF: the value parsing procedure takes care of string, number, and

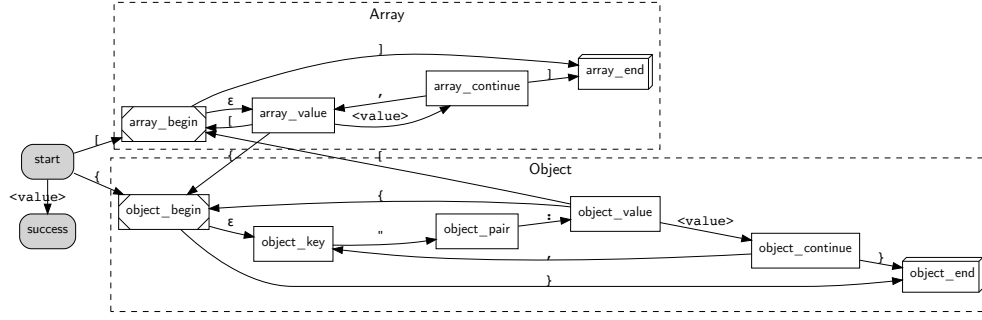


Figure 3: State transition diagram of FSM-based recursive descent parsing of JSON documents. Here “start” is the start state, “succeed” is the only accept state. Nodes with solid diagonal and 3D box markers indicate recursion-entering states and recursion-exiting states, respectively. Edge labels indicate characters that trigger the transition, $\langle \text{value} \rangle$ is shorthand for “ $| \text{t} | \text{f} | \text{n} | - | 0-9$ ”, i.e. all beginning characters of non-structural JSON values, ϵ indicates ϵ -transition.

keyword values; the array parsing procedure merges the nonterminals array and array-elements, and similarly for objects.

3.3.2 Finite State Machine

As mentioned in §2.3, recursive descent parsers do not have to be implemented as recursive procedures. An alternative implementation is a finite state machine with stack, where each state represents a position in one production rule. The state transition diagram for JSON is illustrated in figure 3. To demonstrate how states are derived, we take the production rule for JSON array as an example:

```
array          = "[", (1) array-elements, "]" (4)
array-elements = (2) value, (3) [ ",", (2) array-elements ]
```

The brackets indicate different states derived from the production rule, brackets with the same number share the same state. These states correspond to “array_begin”, “array_value”, “array_continue”, and “array_end” in the diagram.

Ignoring the stack, our FSM is essentially a deterministic finite automaton (DFA). However, DFAs alone cannot describe JSON grammar, since JSON is context-free while DFAs only recognize regular languages. This is where the stack becomes useful. When a recursion-entering state (marked with solid diagonals) is reached, the state before the transition is pushed onto the stack. When a recursion-exiting state (marked with dashed diagonals) is reached, the state from stack top (denoted as S) is popped from stack, and the next state becomes the state that arrives from S by following the $\langle \text{value} \rangle$ -transition.¹

3.3.3 Shift-Reduce

We then moved towards multi-threaded parsing. Among various design choices, we decided to use a fixed coarse-grained work assignment — divide the array of structural characters into chunks, and have each thread construct a partial parse tree for each chunk. This requires the capability of parsing incomplete input into a partial structure.

Our first attempt at multi-threaded parsing was the shift-reduce parser. Since shift-reduce is bottom-up and creates subtrees first, we believe it is naturally suitable for partial parsing.

Shift-reduce parsing also maintains a stack, where each element is a parse tree, i.e., a JSONValue node. We extend JSONValue by introducing three new types: *partial-array*, *partial-object*, and

¹This is because JSON arrays and objects are also considered values. Exiting from \ast_end states denotes finished parsing of a complete JSON object or array, which is equivalent to a parsed value.

```

        "[" , "]" => array
        "[" , value , "]" => array
        "[" , partial-array , "]" => array

        "{" , "}" => object
        "{" , partial-object , "}" => object

        "[" , value , "," => partial-array , ","
partial-array , "," , value , "," => partial-array , ","
        "," , value , "," => "," , partial-array , ","

partial-object , "," , string , ":" , value => partial-object
        string , ":" , value => partial-object

```

Figure 4: Production rules of JSON grammar used in the shift-reduce parser. Quoted character matches structural type nodes with the same character, and “value” matches any fully-parsed JSON value, i.e. anything except *partial-array*, *partial-object*, and *structural*. Note that certain right hand sides contain more than one node, in which case they’re all pushed onto the stack, from left to right.

```

partial-object , "," , partial-object => partial-object
partial-array , "," , partial-array => partial-array

```

Figure 5: Extra production rules used in the merge process of the shift-reduce parser.

structural. A *structural* node consists of a single structural character, while *partial-array* and *partial-object* correspond to the “array-elements” and “object-elements” nonterminals. The algorithm proceeds by alternating between shift and reduce steps. In a shift step, one input character is consumed and a new node is pushed onto the stack. In a reduce step, the stack top is examined against production rules, and when a match is found, replaces the nodes with a joined parse tree. The specific rules are as follows:

Shift If the character is the beginning character of a non-structural JSON value (`"|t|f|n|-|0-9`), parse the value, and push a node of corresponding type onto stack. Otherwise, push a node of type *structural* onto stack.

Reduce Figure 4 lists all the production rules used in the reduce step. The left hand side indicates the pattern to match, with the rightmost element being the stack top; all nodes on the left side are joined as one parse tree with the type specified on the right hand side. For the general shift-reduce parser, an arbitrary number of reduce steps may be performed after each shift step. A nice property here is that there can be at most two reduce steps after each shift step, because all productions rules that yield *array*, *object*, and *partial-array* requires different structural characters at stack top.

After each thread has finished parsing, another round of shift-reduce is performed to merge the stacks from all threads. The same set of production rules are used, with a few extensions as listed in figure 5. Note the number of reduces steps after each shift step is still bounded to 2.

There are two implementation details that are crucial for performance:

1. *partial-array* and *partial-object* nodes are also linked lists. An additional pointer to the end of list must be stored, otherwise appending a value to the partial structure would degenerate to linear complexity.
2. Our set of production rules is carefully crafted to aggressively construct partial structures at the earliest possible. This guarantees that the final stack size is linear to the maximum nested depth of the parsed chunk.

Although seeming promising, the performance of shift-reduce parser (even with multiple threads) is much worse than the FSM method. We observed $5\times$ longer run time for single-threaded shift-reduce. Although running with 4 threads resulted in $2\times$ speedup compared to the single-thread version, it’s still much slower than FSM. Reasons for degeneration are two-fold:

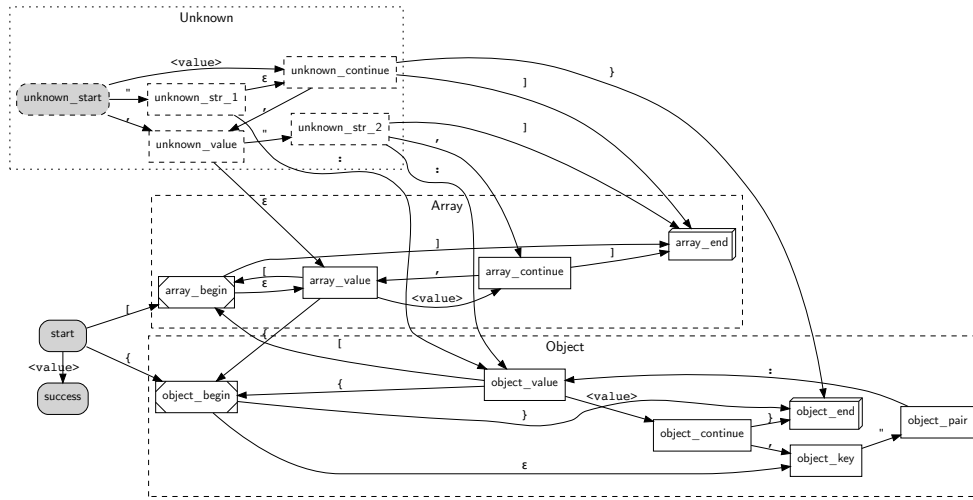


Figure 6: State transition diagram of FSM-based recursive descent parsing of JSON documents with unknown states. Here “unknown_start” is the start state for partial parsing. Other notations are identical to figure 3.

1. Increase in the number of branches per structural character. After each shift step, no less than 4 rules are matched against the stack, which translates to at least 5 branches. For the FSM approach, only one branch is required.
2. Increase in the number of memory accesses. The use of nodes in itself resulted in more memory writes. Moreover, an additional partial node is created for each structure and for each structural character.

3.3.4 Multi-Thread FSM

At this point, we had to go back to recursive descent parsing that we thought was impossible to parallelize, because it’s non-trivial to determine the state of a position in the middle of the input string without actually parsing until that position. However, we realized that this is not true for our case, because we only consider structural characters. For a grammar as simple as the JSON grammar, any state can be easily determined with a few lookahead characters.

With this in mind, we redesigned the FSM to incorporate “unknown states”, as illustrated in figure 6. When a thread starts parsing in the middle of the input string, it begins in the “unknown_start” state. Each unknown state can be considered as a set of states in the original FSM, and our new FSM is essentially the powerset construction of the original FSM treated as an NFA.

An issue that would arise only when parsing a partial structure is: transitioning to a recursion-exiting state with an empty stack. This indicates we have encountered an unmatched closing bracket, so we record the bracket position in a separate list, and transition to the “unknown_continue” state (the state arriving from “unknown_start” by following the <value>-transition).

Similar to multi-threaded shift-reduce, after each thread has finished parsing, we merge the partial structures to form a complete parsed document. Here we only need to match the opening brackets that are left in the FSM stack, with the closing brackets that are stored in the separate list.

However, a pitfall here is that even if each thread contains valid partial JSON and all brackets match, it is still possible for the JSON document to be malformed. For instance, consider the examples in table 1. All examples are malformed JSON documents, when parsed with two threads, each chunk can be accepted by the FSM, and brackets match across chunks. Such grammar violations can only appear at chunk borders, so it suffices to check the characters around borders to rule out these exceptions.

Table 1: Malformed JSON examples. When parsed with two threads, each chunk can be accepted by the FSM, and brackets match across chunks.

Invalid Reason	1st Thread	2nd Thread
Missing colon (:	{ "1": 2, "3"	4, "5": 6 }
Missing comma (,)	[1, 2	3, 4]
Missing comma (,)	[[1, 2]	[3, 4]]
Extra colon (:	{ "1": 2, "3":	: 4, "5": 6 }
Extra comma (,)	[1, 2,	, 3, 4]
Extra kv-pair in object	{ "1": 2, "3":	"4": 5, "5": 6 }
Array value in object	{ "1": 2,	"3", "5": 6 }

This method enjoys the efficiency of FSM and tape-bases storage, while also allowing parallelization. By combining the best of both worlds, we obtain our fastest parsing method.

3.4 Parsing of Values

Another key process in parsing JSON is the parsing of JSON values, especially numbers and strings.

3.4.1 Number Parsing

Number parsing is required to convert a character string into numeric values, which could be either integer or decimal. Built-in functions like `atoi` can handle this, but with the uncertainty of numeric type, this is probably not the optimal choice. In real world datasets, many JSON files contain many numbers. Thus, parsing of numbers can become a bottleneck if not designed efficiently.

Here we first reimplemented the SIMD-optimized version of number parsing as used in `simdjson`. The key idea here is to parse digits in groups of 8, and in each group, iteratively multiply-add adjacent values and merge results in a divide-and-conquer fashion using instructions like `pmaddubsw`. Many floating point numbers have long decimal parts, which can benefit greatly from this optimization. Algorithm 1 describes this process.

Algorithm 1 Number Parsing

```

1: procedure PARSENUMBER(input)
2:    $s \leftarrow$  current parsing position
3:   Check if number is negative
4:   Convert digits one by one for integer part
5:   if  $s = \text{"."}$  then
6:     if next 8 characters are digits then
7:       ParseNumberSIMD( $s + 1$ )
8:     Convert remaining digits one by one
9:   if  $s = \text{"e"}$  or  $s = \text{"E"}$  then
10:    Check if exponent is negative
11:    Convert digits one by one for exponent
12:    Multiply number by precomputed exponent of 10
13:  return parsed number
14: procedure PARSENUMBERSIMD(input)
15:  Load a group of 8 digit characters
16:  Subtract digit characters with character "0"
17:  while Not fully reduced do
18:    Multiply position factor and add adjacent elements
19:  return parsed number

```

3.4.2 String Parsing

String parsing can be split into two separate tasks:

- Relocation (memory move). The relocation of string values has both pros and cons. If strings are relocated to another memory position, we can discard the input string which can save space and make the memory organization of the JSON document more compact. The cost is more memory write during parsing which may incur higher cache miss rate and increase peak memory usage.
- Unescape values. The original string contains backslashes to escape their following characters. These characters need to be unescaped to their original values in this process. We could reuse procedures in stage 1 to locate escape positions, but we still have to deal with conversion and position shifts (backslashes need to be overwritten).

We implemented several different methods for string parsing:

- `simdjson` used a simple string parsing method. After retrieving the escape mask, we can use the `tzcnt` instruction to quickly find the position of the first escape character. We then convert the character and shift the following characters one position ahead, then perform the process again starting from the next position. This process is performed as many times as the number of escaped characters, which is usually small for real world JSON documents.
- One slight improvement based on the `simdjson` method is to convert every escaped character in the current vector before moving to the next. By using the `blzr` instruction, we can erase the last non-zero bit in the escape mask and then continuing using `tzcnt` to find the next escape position. In our experiments, we found this method to be slightly more efficient. Algorithm 2 describes this process.
- Another attempt we made to improve string parsing is to fully vectorize this whole process. While it is rather convenient to use the `vcompress` instruction in AVX512, it is much harder for ordinary machines to do so with AVX2 support only. Using only AVX2 resulted in the fully-vectorized version have a large overhead. In tests using artificial data with about 6% of characters being escaped, this method can be 2 to 3 times faster than the previous one. However, real world JSON documents contain much fewer than 6% escaped characters, and as a result, its performance is worse than the previous one on real datasets. The idea is an extension of a vectorized space eliminator based on the `pext` instruction proposed in [5]. Algorithm 3 describes this process.

Algorithm 2 String Parsing

```
1: procedure PARSESTRING(input)
2:   Divide input into groups
3:   for each group do
4:     Compute the escape mask of the group
5:     while mask is not empty do
6:       Find position of next escaped character
7:       Convert escaped character and move data
8:       Erase last non-zero bit in mask
9:   return
```

Algorithm 3 String Parsing Fully Vectorized

```
1: procedure PARSESTRINGFULLYVECTORIZED(input)
2:   Divide input into groups
3:   for each group G do
4:     Extract the escape mask M of the group
5:     Transform all characters into E using two way vpshufd
6:     Merge E and G with mask M using vpblendvb
7:     Extract each bit position's de-escaped bit using pext
8:     Reconstruct result with bits extracted above
9:   return
```

Table 2: Dataset statistics. The last column contains the structural character count. Datasets in bold are large datasets we added.

Dataset	Bytes	Integer	Decimal	String	Object	Array	null	true	false	Structural
apache_builds	127275	2	0	5289	884	3	0	2	1	12365
canada	2251051	46	111080	12	4	56045	0	0	0	334374
citm_catalog	1727204	14392	0	26604	10937	10451	1263	0	0	135991
github_events	65132	149	0	1891	180	19	24	57	7	4657
gsoc-2018	3327831	0	0	34128	3793	0	0	0	0	75842
instruments	220346	4935	0	6889	1012	194	431	17	109	27174
lottery	307788	5350	0	12122	114	1791	4439	2	6	47656
marine_ik	2983466	130225	114950	38268	9680	28377	0	6	0	643013
mesh	723597	40613	32400	11	3	3610	0	0	0	153275
mesh.pretty	1577353	40613	32400	11	3	3610	0	0	0	153275
numbers	150124	0	10001	0	0	1	0	0	0	20004
random	510476	5002	0	33005	4001	1001	0	495	505	88018
twitter	631514	2108	1	18099	1264	1050	1946	345	2446	55264
twitterescaped	562408	2108	1	18099	1264	1050	1946	345	2446	55164
update-center	533178	0	0	27229	1896	1937	0	134	252	63420
citylots	189778220	0	7875189	5109655	619675	3038859	54336	0	0	33395428
crime	1041644856	7838495	0	50157642	620	3919022	14507113	2	1959464	156764725
motor	579937595	4963886	0	34881548	394	1654520	4827687	4	6	92656098
traffic	1038464885	4541077	0	67473705	574	3027295	8210214	2	1513625	169532998

Table 3: System, compiler, and hardware information of each testing machine.

System	Compiler	CPU	RAM
Windows 10 Subsystem for Linux	GCC 7.3.0	Intel i7-8805H 2.6GHz (Coffee Lake)	32 GB
Ubuntu 18.04.2 LTS	GCC 7.3.0	Intel i5-6200U 2.3GHz (Skylake)	8 GB
macOS Mojave 10.14.4	Clang 10.0.1	Intel i7-4870HQ 2.5GHz (Haswell)	16 GB

3.4.3 Keyword Parsing

The keyword values `true`, `false`, `null` are also required to parse. Here, since these values are all smaller than 64 bits, parsing could be implemented by comparison of whole 64-bit words to accelerate.

3.4.4 Multi-Threaded Value Parsing

Since string values and numeric values constitute a large part of many JSON documents, we observe that the parsing function is time consuming so we set up dedicated threads to do parsing of these values. For example, each separate string parsing thread will scan a range of structural characters, and perform string parsing when a quote character is encountered. Number parsing threads perform similarly.

4 Experiments

To test the performance of our algorithm, we conducted a series of experiments to compare with baseline algorithms including *simdjson* and *rapidjson*.

Our test data include the benchmark tests used in *simdjson*, and several very large JSON documents retrieved from the U.S. government open data website. Detailed statistics of each test case is listed in table 2. The additionally added large datasets include *citylots*, *crime*, *motor*, and *traffic*. Although *simdjson* claims to reach the speed of 2 GB/s, their performance was evaluated on small JSON documents of several hundred kilobytes, largest of which did not exceed 3 MB. Although still impressive, we believe these documents alone does not fully describe the performance characteristics of parsing algorithm, which is why we included real world large documents of several hundred megabytes or even reaching gigabytes. On these documents, *simdjson* performance drops to around 400 MB/s.




Figure 7: Cycles per input byte on Intel i5-6200U 2.3GHz CPU (Skylake). Here the MercuryParser is using only 4 state machine transition threads without dedicated string parsing threads. If using dedicated string parsing threads, cycle count in stage 2 will be extremely small since the main thread does not perform string parsing.

Table 4: Parsing performance (MB/s) on Intel i7-4870HQ 2.5GHz CPU (Haswell). Each result listed is the minimum running time over 1000 executions for small cases and 10 executions for large cases.

Dataset	MercuryJson		simdjson	rapidjson
	default	single-thread		
apache_builds	888.83	1016.71	1846.84	315.63
canada	1396.96	657.01	900.76	310.18
citm_catalog	2415.21	1823.48	2240.84	657.40
github_events	538.64	836.27	2029.54	339.89
gsoc-2018	3128.66	2806.44	2864.82	397.88
instruments	1216.33	1238.95	1699.19	401.69
lottery	1232.74	1110.72	1403.55	304.53
marine_ik	1134.65	607.11	763.65	246.77
mesh	971.40	531.70	812.59	307.49
mesh.pretty	1559.63	949.90	1235.20	543.38
numbers	773.10	556.50	963.35	313.44
random	1356.10	1236.63	1210.76	239.67
twitter	1895.65	1650.43	1916.82	367.24
twitterescaped	1453.29	1353.39	1070.30	296.34
update-center	1509.47	1350.83	1569.28	253.60
citylots	558.16	374.80	457.43	231.78
crime	671.21	528.27	480.87	225.71
motor	623.72	513.35	498.39	202.90
traffic	644.75	508.32	500.92	198.54

4.1 Experiment Settings

We collect the best performance over 1000 executions for small datasets and 10 executions for large datasets. When measuring run time, we follow *simdjson* settings and do not take into account the time required for disk I/O and memory allocation.

All libraries and testing programs are compiled with optimization level 3 (-O3) and targets native processor architecture only (-march=native). We test all programs on three different machines, details of which are listed in table 3.

MercuryJson The default configuration of our method is using multi-threaded FSM with 4 threads, and 4 other dedicated threads to parse string values. We also test our performance under a “single-threaded” scenario, where we use the single-threaded FSM method, but still using 4 dedicated string parsing threads. Note that the single-threaded FSM is optimized and has fewer branches per structural character than the multi-threaded implementation.

simdjson We use the bundled benchmark script (benchmark/parse.cpp) in the *simdjson* package for testing. We modify the execution counts² to match our testing scheme. We compile the script using the default preprocessor flags.

rapidjson *rapidjson* is a header-only library. We run parsing under “in-situ” (in-place) mode for fairer comparison.

²By default, *simdjson* executions 10 times for documents over 1 MB.

Table 5: Parsing performance (MB/s) on Intel i7-8850H 2.6GHz CPU (Coffee Lake). Each result listed is the minimum running time over 1000 executions for small cases and 10 executions for large cases. This program is executed in the Windows Subsystem for Linux which has potential performance issues. A more accurate benchmark is in table 4 which is executed on UNIX-based macOS.

Dataset	MercuryJson		simdjson	rapidjson
	default	single-thread		
apache_builds	384.26	777.07	2043.94	772.13
canada	1691.97	605.49	1045.25	627.51
citm_catalog	2294.46	2585.45	2637.36	1092.88
github_events	195.58	422.26	1616.18	807.73
gsoc-2018	3233.49	3234.48	2586.53	992.67
instruments	383.39	532.81	1812.06	859.46
lottery	406.61	1232.28	1617.38	716.96
marine_ik	1499.32	535.29	969.79	591.62
mesh	1067.07	813.10	986.37	562.73
mesh.pretty	738.51	1367.28	1409.11	873.97
numbers	258.29	655.24	1057.96	295.87
random	411.87	1371.35	1465.20	591.31
twitter	1094.42	1817.87	2202.70	805.37
twitterescaped	515.73	1444.53	1214.97	695.93
update-center	686.30	631.02	1846.82	643.40
citylots	303.74	335.83	391.65	557.97
crime	371.71	456.30	385.63	658.60
motor	350.71	477.57	401.62	614.86
traffic	298.79	424.54	390.84	616.43

4.2 Results

The detailed results of all experiments are listed in tables 4, 5, and 6. We compare between *simdjson*, *rapidjson*, and two different settings of our method (MercuryJson). Results in bold indicates that our method outperforms the state-of-the-art method *simdjson*.

4.3 Analysis

In table 8, we show performance details obtained from hardware counters when parsing *citylots* with Intel i5-6200U 2.3GHz CPU (Skylake). Measurements in stage 1 are similar between MercuryJson and *simdjson*. In stage 2, multi-threading greatly helps the program reduce the number of instructions executed on a single thread, and as a result, branch misses and cache misses are greatly reduced as well, which yields a much faster processing speed (cycles/byte) in spite of a lower instruction processing rate (ins/cycles).

There are several factors that could make differences to algorithm performance:

Document Size The size of documents significantly affects the performance. This is mainly due to the overhead brought by threads. In stage 2, the average run time per structural character is only several cycles. Launching 4 string parsing threads takes about 300,000 cycles and joining them takes 200,000 cycles. Launch 4 FSM threads takes another 500,000 cycles. Compare this to the sequential algorithm, where each structural character only takes 15 to 30 cycles on average, including the parsing of values. As a result, the overhead of our multi-threaded algorithm dominates the run time on documents such as *github_events* (4657) or *apache_builds* (12365) which only

Table 6: Parsing performance (MB/s) on Intel i5-6200U 2.3GHz CPU (Skylake). Each result listed is the minimum running time over 1000 executions for small cases and 10 executions for large cases. This machine has only 8 GB RAM, *simdjson* experiences out-of-memory errors for most of the large documents so the performances for those cases are not available.

Dataset	MercuryJson		simdjson	rapidjson
	default	single-thread		
apache_builds	721.12	878.52	1630.73	541.88
canada	741.23	513.20	754.91	415.99
citm_catalog	1312.04	1471.27	1815.90	765.12
github_events	490.35	791.75	1742.76	562.00
gsoc-2018	1985.50	1974.01	2142.33	662.57
instruments	804.36	1037.75	1463.42	588.05
lottery	676.55	885.37	1163.98	491.60
marine_ik	603.10	446.00	664.64	393.88
mesh	611.55	620.24	429.66	377.59
mesh.pretty	1021.44	772.40	603.42	576.67
numbers	653.91	557.57	785.20	331.79
random	658.56	806.83	1046.91	403.62
twitter	1079.58	1277.35	1611.96	563.41
twitterescaped	876.13	910.83	841.35	449.40
update-center	904.28	872.28	1307.85	436.69
citylots	558.16	451.11	510.76	353.05
crime	625.25	619.76	–	421.04
motor	623.57	623.27	–	379.82
traffic	594.75	612.15	–	374.45

Table 7: Time proportion for two stages in MercuryJson with standard configuration

Dataset	Stage 1	Stage 2
apache_builds	7.58%	92.42%
canada	36.04%	63.96%
citm_catalog	39.76%	60.24%
github_events	3.67%	96.33%
gsoc-2018	49.05%	50.95%
instruments	7.24%	92.76%
lottery	8.02%	91.98%
marine_ik	33.51%	66.49%
mesh	18.64%	81.36%
mesh.pretty	13.40%	86.60%
numbers	6.06%	93.94%
random	8.61%	91.39%
twitterescaped	10.99%	89.01%
twitter	25.22%	74.78%
update-center	16.51%	83.49%
citylots	40.67%	59.33%
crime	46.27%	53.73%
motor	43.90%	56.10%
traffic	39.76%	60.24%

Table 8: Performance details on citylots with Intel i5-6200U 2.3GHz CPU (Skylake). Stage 2 statistics for MercuryJson evaluated on a single thread due to measurement limitations.

Stage	Parser	instructions	cycles	ins/cycles	mis. branches	cache access	cache miss	cycles/byte
1	MercuryJson	584761388	228997694	2.55	418985	6549567	4848937	1.21
	simdjson	656284673	225680027	2.91	408001	6482802	4805979	1.19
2	MercuryJson	364347197	214925286	1.70	188544	5491162	2806829	1.13
	simdjson	1660408275	518461548	3.20	758999	12624901	8210427	2.73

have thousands of structural characters to process in stage 2. For large JSON documents cases such as `crime` (1.57×10^8) and `traffic` (1.70×10^8), this overhead can be ignored.

Instruction Count The implementation of `simdjson` is optimized to extreme and there is hardly any redundant instructions. However, our implement of the multi-thread FSM requires executing more instructions than a single-threaded version. For example, we need to check whether the stack is empty when transitioned to a recursion-exiting state, and we need to write extra information on tape for use in the merging process. This increases the instruction count in stage 2.

Cache Access & Misses Just as the increase in instruction count matters, cache access is also a key factor. Decrease in cache miss rate can greatly benefit performance. Both stages of parsing have compute-intensive workload and possess strong data dependencies between characters, which means the cache miss penalty cannot be effectively alleviated by out-of-order execution and other internal optimization methods.

Branch Misses Branch miss rate is also important for similar reasons. During branch misprediction, partially executed instructions will be discarded and the instruction pipeline restarts, incurring a delay. Branches are frequent in the instruction stream since there are many validations and state transitions in stage 2, so carefully designed code with a lower branch miss rate will be crucial to the overall performance.

Context Switching This is another overhead brought by multi-threading. With the use of multiple threads, operation system could make context switches between threads which will harm the performance. Although this problem could be minimized by assigning suitable numbers of threads, due to the potential work imbalance, this problem could still exists as we observed an increase number of context switches and page fault events using `perf`.

Platform We conducted our experiments on three different platforms including WSL (Windows Subsystem for Linux), Ubuntu 18.04, and macOS Mojave 10.14.4. The performance on WSL seems to be not very reliable where `rapidjson` runs the fastest on large data which is out of our expectation. The behaviour on the other native UNIX based platform is similar and close to our expected results.

Number of Threads Figure 8 shows the parsing speed when using different numbers of FSM threads. Experiments are performed on a 4 core CPU. Since JSON parsing is computation bounded, there are no significant improvements when using more than 4 threads. Significant improvement can be observed from using 1 thread to 4 threads. Also, the use of dedicated string parsing threads can also boost performance as well, as long as the total number of threads does not exceed 4.

5 Conclusion

In conclusion, using multi-threading to optimize JSON parsing is possible and could be useful for large JSON files. The overhead of multi-threading makes it hard to compete with single thread based parsing algorithm on small cases. The SIMD optimization, on the other hand, can accelerate the parsing under different scenarios which makes its performance more robust. There is still some room for further optimization. With our multi-threading optimization, the second stage of parsing is not the bottleneck any more and the first stage currently is not optimized with multi-threading.

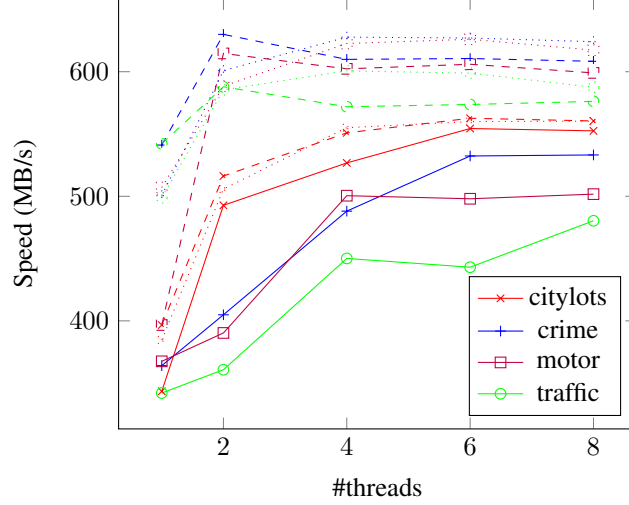


Figure 8: Speed up with different number of threads for FSM parser. Solid line: no dedicated string parsing threads. Dashed line: 2 dedicated string parsing threads. Dotted line: 4 dedicated string parsing threads. Experiments executed on Intel i5-6200U 2.3GHz CPU (Skylake) with 4 cores.

It might be not trivial to develop a multi-thread version for that as the parsing in stage 1 always depends on limited mask results from previous step. We expect performance could be improved further on top of that.

6 Division of Work

Da reimplemented stage 1 as described in `simdjson` (extracting structural characters using SIMD). Zecong and Da collaboratively worked on the re-implementation of the tape-based storage and FSM parser in stage 2. Zecong implemented the shift-reduce parser and multi-threading for FSM parser. Both students performed benchmark testing.

References

- [1] Tim Bray. The javascript object notation (json) data interchange format. Technical report, 2017.
- [2] W.H. Burge. *Recursive programming techniques*. Addison-Wesley Series in Electrical and Computer Engineering. Addison-Wesley Longman, Incorporated, 1975.
- [3] Douglas Crockford. The application/json media type for javascript object notation (json). Technical report, 2006.
- [4] Geoff Langdale and Daniel Lemire. Parsing gigabytes of JSON per second. *arXiv preprint arXiv:1902.08318*, 2019.
- [5] Daniel Lemire. Despace. <https://github.com/lemire/despacer>, 2017.
- [6] Yanan Li, Nikos R Katsipoulakis, Badrish Chandramouli, Jonathan Goldstein, and Donald Kossmann. Mison: a fast json parser for data analytics. *Proceedings of the VLDB Endowment*, 10(10):1118–1129, 2017.